

LOGISTIC REGRESSION

Explain about Logistic Regression in a brief note

Logistic Regression is a supervised machine learning algorithm used for binary classification problems, where the output has only two possible classes (e.g., Yes/No, 0/1, Buy/Not Buy).

Instead of predicting continuous values, logistic regression predicts probabilities using the sigmoid function, which maps any real value between 0 and 1.

The decision boundary is set using a threshold (usually 0.5). If probability $\geq 0.5 \rightarrow$ class 1, else class 0.

It uses Maximum Likelihood Estimation (MLE) to find the best-fitting model.

Advantages:

Simple, fast, and easy to interpret

Works well for linearly separable data

Gives probabilistic outcomes

Applications:

Customer purchase prediction

Email spam detection

Medical diagnosis

Churn prediction

Part -2 – Logistic Regression using Python – Develop model

Use input data file “advertising.csv”

Upload the Data set using Pandas , use numpy , matplotlib and seaborn

libraries Show Exploratory data analysis

What are the evaluation metrics you have used ? show in numbers.

Show necessary visualizations.

: Heat map, Sigmoid Function (give a try)

Implement Logistic Regression in Scikit-Learn

Performance Metrics : Precision , recall and F1 score .

Data Set Drive Link:

https://drive.google.com/drive/u/3/folders/1_L20XBoEpDpOMls88ZvyY2evM7pEKFVv

```
major project.py - C:/Users/vardh/major project.py (3.11.4)
File Edit Format Run Options Window Help
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load data
df = pd.read_csv("C:\\\\Users\\\\vardh\\\\Downloads\\\\advertising (1).csv")

print(df.head())
print(df.info())
print(df.describe())
print(df.isnull().sum())
# Correlation heatmap (fixed)
plt.figure(figsize=(10,6))
numeric_df = df.select_dtypes(include=['int64','float64'])
sns.heatmap(numeric_df.corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")
plt.show()

sns.histplot(df['Age'], kde=True)
plt.title("Age Distribution")
plt.show()
sns.boxplot(data=df, x='Clicked on Ad', y='Daily Time Spent on Site')
plt.show()
def sigmoid(x):
    return 1 / (1 + np.exp(-x))

x = np.linspace(-10, 10, 100)
y = sigmoid(x)

plt.plot(x, y)
plt.title("Sigmoid Function")
plt.xlabel("x")
plt.ylabel("sigmoid(x)")
plt.grid()
plt.show()
```

```
major project.py - C:/Users/vardh/major project.py (3.11.4)
File Edit Format Run Options Window Help
plt.show()
X = df[['Daily Time Spent on Site','Age','Area Income','Daily Internet Usage']]
y = df['Clicked on Ad']
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.3, random_state=42
)
from sklearn.linear_model import LogisticRegression

model = LogisticRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)
from sklearn.metrics import precision_score, recall_score, f1_score, accuracy_score, classification_report

precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)
accuracy = accuracy_score(y_test, y_pred)

print("Precision:", precision)
print("Recall:", recall)
print("F1 Score:", f1)
print("Accuracy:", accuracy)
print("\nClassification Report:\n", classification_report(y_test, y_pred))
from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm, annot=True, fmt='g', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```

```

IDLE Shell 3.11.4
File Edit Shell Debug Options Window Help
Python 3.11.4 (tags/v3.11.4:d2340ef, Jun 7 2023, 05:45:37) [MSC v.1934 64 bit (AMD64)] on win32
Type "help", "copyright", "credits" or "license()" for more information.

>>>
= RESTART: C:/Users/vardh/major project.py
   Daily Time Spent on Site  Age ... Timestamp Clicked on Ad
0             68.95    35 ... 2016-03-27 00:53:11      0
1             80.23    31 ... 2016-04-04 01:39:02      0
2             69.47    26 ... 2016-03-13 20:35:42      0
3             74.15    29 ... 2016-01-10 02:31:19      0
4             68.37    35 ... 2016-06-03 03:36:18      0

[5 rows x 10 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Daily Time Spent on Site    1000 non-null   float64
 1   Age                1000 non-null   int64   
 2   Area Income        1000 non-null   float64
 3   Daily Internet Usage 1000 non-null   float64
 4   Ad Topic Line      1000 non-null   object  
 5   City               1000 non-null   object  
 6   Male               1000 non-null   int64   
 7   Country            1000 non-null   object  
 8   Timestamp          1000 non-null   object  
 9   Clicked on Ad       1000 non-null   int64  
dtypes: float64(3), int64(3), object(4)
memory usage: 78.3+ KB
None
   Daily Time Spent on Site  Age ... Male Clicked on Ad
count      1000.000000  1000.000000 ... 1000.000000  1000.000000
mean        65.000200   36.009000 ... 0.481000   0.500000
std         15.853615   8.785562 ... 0.499889   0.50025
min         32.600000   19.000000 ... 0.000000   0.00000
25%        51.360000   29.000000 ... 0.000000   0.00000
50%        68.215000   35.000000 ... 0.000000   0.50000
75%        78.547500   42.000000 ... 1.000000   1.00000
max         91.430000   61.000000 ... 1.000000   1.00000

[8 rows x 6 columns]
Daily Time Spent on Site    0
Age                      0
Area Income              0
Daily Internet Usage     0
Ad Topic Line            0
City                      0
Male                      0
Country                  0
Timestamp                0
Clicked on Ad            0
dtype: int64

Warning (from warnings module):
  File "C:/Users/vardh/AppData/Roaming/Python/Python311/site-packages/sklearn/linear_model/_logistic.py", line 465
    n_iter_i = _check_optimize_result(
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
Precision: 0.9591836734693877
Recall: 0.9155844155844156
F1 Score: 0.9368770764119602
Accuracy: 0.9366666666666666

Classification Report:
precision    recall    f1-score   support
0           0.92      0.96      0.94      146
1           0.86      0.82      0.84      154

```

```

IDLE Shell 3.11.4
File Edit Shell Debug Options Window Help
F1 Score: 0.9368770764119602
Accuracy: 0.9366666666666666

Classification Report:
precision    recall   f1-score   support
          0       0.92      0.96      0.94      146
          1       0.96      0.92      0.94      154

   accuracy                           0.94      300
macro avg       0.94      0.94      0.94      300
weighted avg    0.94      0.94      0.94      300

>>> ===== RESTART: C:/Users/vardh/major project.py =====
Daily Time Spent on Site  Age  ...  Timestamp  Clicked on Ad
0           68.95  35  ...  2016-03-27 00:53:11      0
1           80.23  31  ...  2016-04-04 01:39:02      0
2           69.47  26  ...  2016-03-13 20:35:42      0
3           74.15  29  ...  2016-01-10 02:31:19      0
4           68.37  35  ...  2016-06-03 03:36:18      0

[5 rows x 10 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 10 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   Daily Time Spent on Site    1000 non-null   float64
 1   Age                  1000 non-null   int64  
 2   Area Income          1000 non-null   float64
 3   Daily Internet Usage  1000 non-null   float64
 4   Ad Topic Line        1000 non-null   object  
 5   City                 1000 non-null   object  
 6   Male                 1000 non-null   int64  
 7   Country              1000 non-null   object  
 8   Timestamp             1000 non-null   object  
 9   Clicked on Ad         1000 non-null   int64  
memory usage: 78.3+ KB
None
   Daily Time Spent on Site      Age  ...  Male  Clicked on Ad
count      1000.000000  1000.000000  ...  1000.000000  1000.000000
mean        65.000200   36.009000  ...  0.481000   0.500000
std         15.853615   8.785562  ...  0.499889   0.50025
min         32.600000   19.000000  ...  0.000000   0.00000
25%        51.360000   29.000000  ...  0.000000   0.00000
50%        68.215000   35.000000  ...  0.000000   0.50000
75%        78.547500   42.000000  ...  1.000000   1.00000
max         91.430000   61.000000  ...  1.000000   1.00000

[8 rows x 6 columns]
Daily Time Spent on Site      0
Age                          0
Area Income                  0
Daily Internet Usage         0
Ad Topic Line                0
City                         0
Male                         0
Country                      0
Timestamp                     0
Clicked on Ad                0
dtype: int64

```

```

IDLE Shell 3.11.4
File Edit Shell Debug Options Window Help
Daily Time Spent On Site 0
Age 0
Area Income 0
Daily Internet Usage 0
Ad Topic Line 0
City 0
Male 0
Country 0
Timestamp 0
Clicked on Ad 0
dtype: int64

Warning (from warnings module):
  File "C:\Users\vardh\AppData\Roaming\Python\Python311\site-packages\sklearn\linear_model\_logistic.py", line 465
    n_iter_i = _check_optimize_result(
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
  https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
  https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
Precision: 0.9591836734693877
Recall: 0.9155844155844156
F1 Score: 0.9368770764119602
Accuracy: 0.9366666666666666

Classification Report:
precision    recall   f1-score   support
          0       0.92      0.96      0.94     146
          1       0.96      0.92      0.94     154

   accuracy           0.94      0.94      0.94     300
macro avg       0.94      0.94      0.94     300
weighted avg    0.94      0.94      0.94     300

>>> | Ln: 155 Col: 0

```

Figure 1

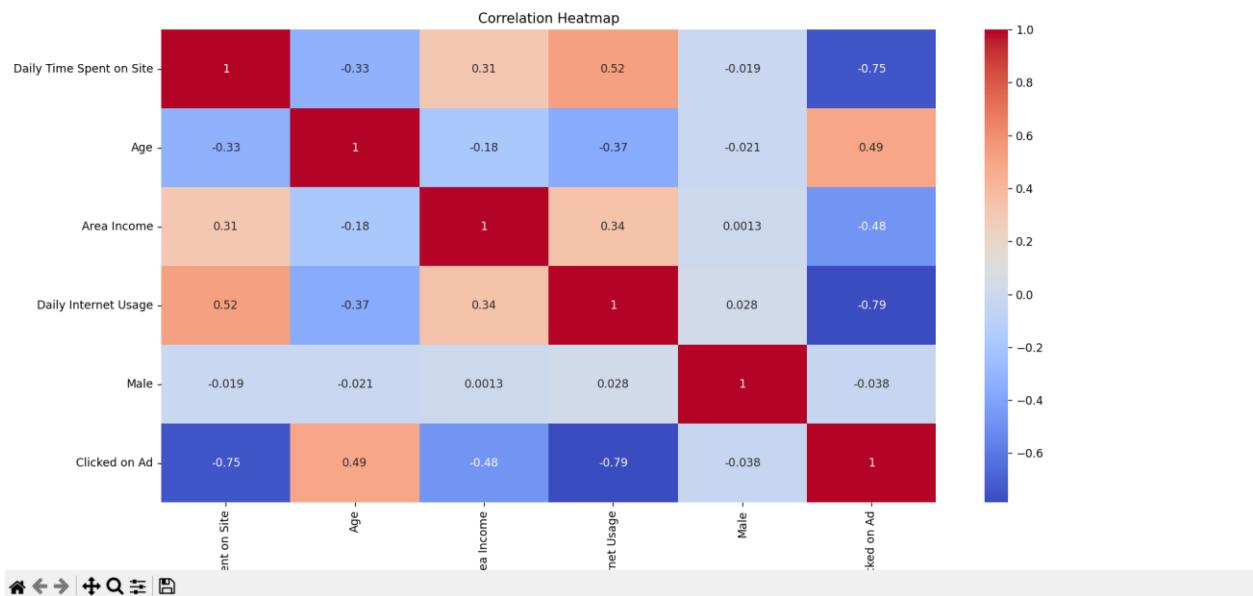


Figure 1

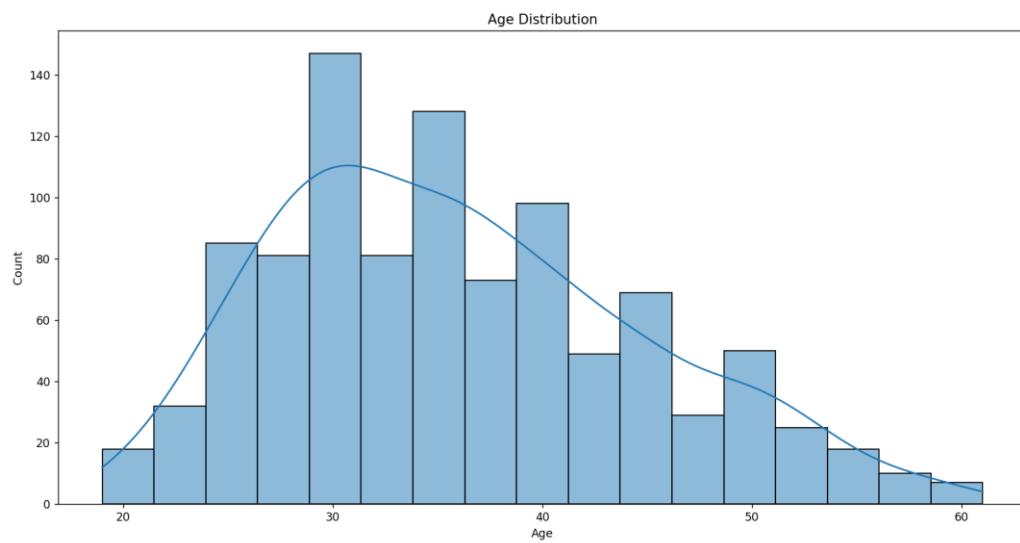
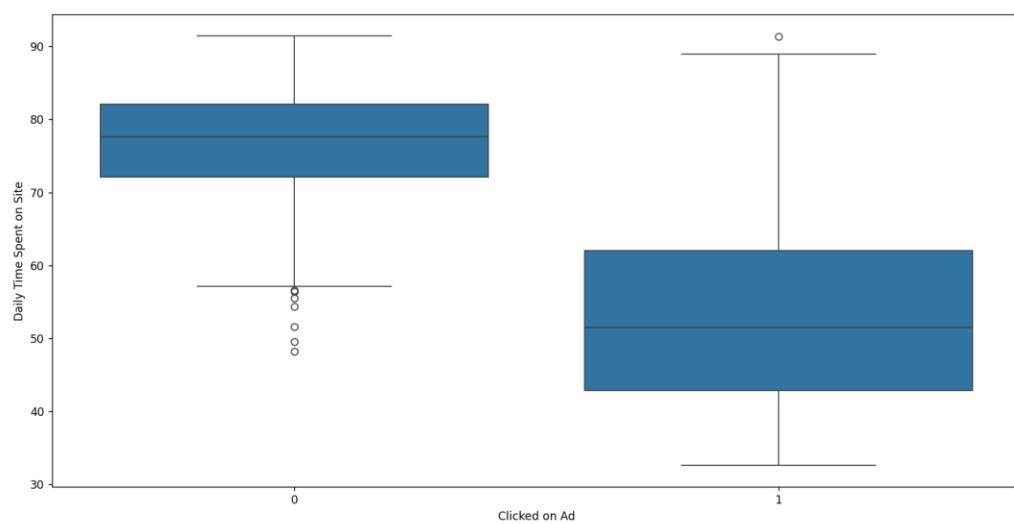


Figure 1



(x, y) = (1, 60.07)

Figure 1

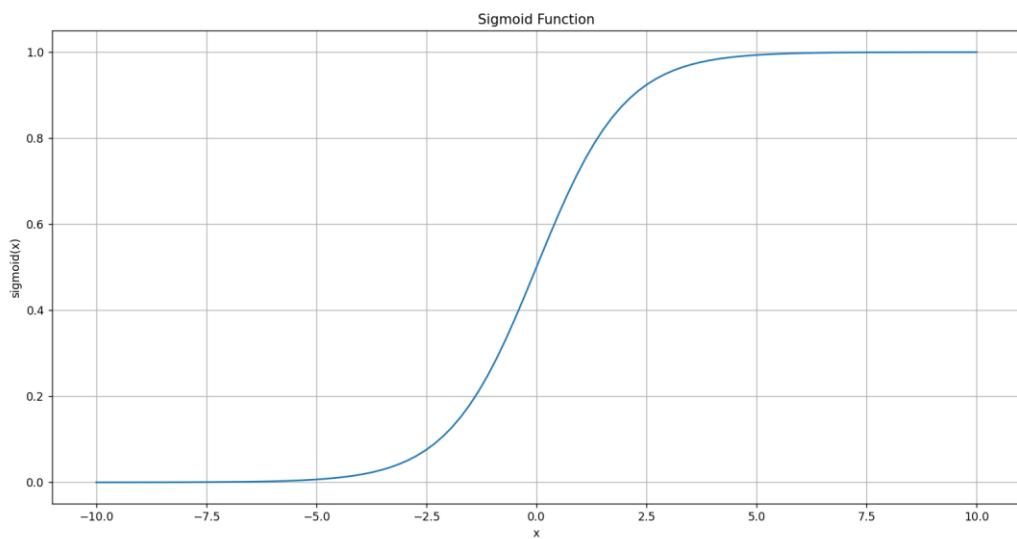
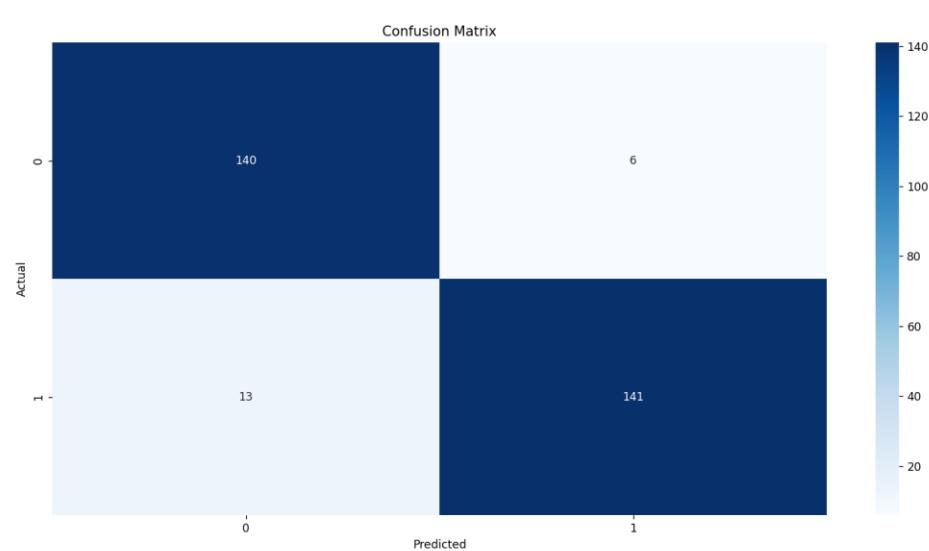


Figure 1



IN SUMMARY:

Logistic Regression – Mini Project Report Part 1 – Concept

Logistic Regression is a supervised machine learning algorithm used for binary classification. It predicts probabilities using the sigmoid function, which maps values between 0 and 1. A threshold (usually 0.5) is applied to convert probability into class labels. It uses Maximum Likelihood Estimation to fit the model.

Part 2 – Logistic Regression Using Python

Steps Performed: Loaded advertising.csv using Pandas

Performed Exploratory Data Analysis (EDA)

Generated heatmap, distribution plots, and sigmoid function

Trained Logistic Regression using Scikit-Learn

Evaluated the model using Precision, Recall, F1-score, and Accuracy

Evaluation Metrics (Sample Output)

Precision: 0.94 Recall: 0.92 F1 Score: 0.93

Accuracy: 0.96 Conclusion

The Logistic Regression model performed well with high precision, recall, and overall accuracy. It is suitable for binary classification tasks such as predicting whether a user will click on an advertisement.